

Test Suite Generation with the Many Independent Objective (MIO) Algorithm

Andrea Arcuri^{a,b}

^a*Westerdals Oslo ACT, Faculty of Technology, Oslo, Norway*

^b*University of Luxembourg, Luxembourg*

Abstract

Context: Automatically generating test suites is intrinsically a multi-objective problem, as any of the testing targets (e.g, statements to execute or mutants to kill) is an objective on its own. Test suite generation has peculiarities that are quite different from other more regular optimisation problems. For example, given an existing test suite, one can add more tests to cover the remaining objectives. One would like the smallest number of small tests to cover as many objectives as possible, but that is a secondary goal compared to covering those targets in the first place. Furthermore, the amount of objectives in software testing can quickly become unmanageable, in the order of (tens/hundreds of) thousands, especially for system testing of industrial size systems.

Objective: To overcome these issues, different techniques have been proposed, like for example the Whole Test Suite (WTS) approach and the Many-Objective Sorting Algorithm (MOSA). However, those techniques might not scale well to very large numbers of objectives and limited search budgets (a typical case in system testing). In this paper, we propose a novel algorithm, called Many Independent Objective (MIO) algorithm. This algorithm is designed and tailored based on the specific properties of test suite generation.

Method: An empirical study was carried out for test suite generation on a series of artificial examples and seven RESTful API web services. The EvOMASTER system test generation tool was used, where MIO, MOSA, WTS and random search were compared.

Results: The presented MIO algorithm resulted having the best overall performance, but was not the best on all problems.

Conclusion: The novel presented MIO algorithm is a step forward in the automation of test suite generation for system testing. However, there

33 are still properties of system testing that can be exploited to achieve even
34 better results.

35 *Keywords:*

36 test generation, SBSE, SBST, multi-objective optimization, system testing

37 1. Introduction

38 Test case generation can be modelled as an optimisation problem, and
39 so different kinds of search algorithms can be used to address it [1]. There
40 can be different objectives to optimise, like for example branch coverage or
41 the detection of mutants in the system under test (SUT). When aiming at
42 maximising these metrics, often the sought solutions are not single test cases,
43 as a single test cannot cover all the objectives in the SUT. Often, the final
44 solutions are sets of test cases, usually referred as *test suites*.

45 There are many different kinds of search algorithms that can be used for
46 generating test suites. The most famous is perhaps the Genetic Algorithms
47 (GA), which is often the first choice when addressing a new software engineer-
48 ing problem for the first time. But it can well happen that on specific problems
49 other search algorithms could be better. Therefore, when investigating a new
50 problem, it is not uncommon to evaluate and compare different algorithms.
51 On average, no search algorithm can be best on all possible problems [2].
52 It is not uncommon that, even on non-trivial tasks, simpler algorithms like
53 (1+1) Evolutionary Algorithm (EA) or Hill Climbing (HC) can give better
54 results than GA (e.g., as in [3]).

55 A major factor affecting the performance of a search algorithm is the
56 so called *search budget*, i.e., for how long the search can be run, usually
57 the longer the better. But the search budget is also strongly related to the
58 tradeoff between the *exploitation* and *exploration* of the search landscape. If
59 the budget is low, then a population-based algorithm like GA (which puts
60 more emphasis on the exploration) is likely to perform worse than a single,
61 more focused individual-based algorithm like HC or (1+1) EA. On the other
62 hand, if the search budget is large enough, the exploration made by the
63 GA can help it to escape from the so called local optima in which HC and
64 (1+1) EA can easily get stucked in.

65 To obtain even better results, then one has to design specialised search
66 algorithms that try to exploit the specific properties of the addressed problem
67 domain. In the case of test suite generation, there are at least the following

68 peculiarities:

- 69 • testing targets can be sought *independently*. Given an existing test suite,
70 to cover the remaining testing targets (e.g., lines and branches), you
71 can create and add new tests without the need to modify the existing
72 ones in the suite. At the end of the search, one wants a minimised test
73 suite, but that is a secondary objective compared to code coverage.
- 74 • testing targets can be strongly related (e.g., two nested branches),
75 as well as being completely independent (e.g., code in two different
76 top-level functions with no shared state).
- 77 • some testing targets can be *infeasible*, i.e., impossible to cover. There
78 can be different reasons for it, e.g., dead code, defensive programming
79 or the testing tool not handling all kinds of SUT inputs (e.g., files or
80 network connections). Detecting whether a target is feasible or not is
81 an undecidable problem.
- 82 • for non-trivial software, there can be a very large number of objectives.
83 This is specially true not only for system-level testing, but also for
84 unit testing when mutation score is one of the coverage criteria [4].
85 Traditional multi-objective algorithms are ill suited to tackle large
86 numbers of objectives [5].

87 In this paper, we propose a novel search algorithm that exploits such
88 characteristics and, as such, it is specialised for test suite generation (and any
89 problem sharing those properties). We call it the Many Independent Objective
90 (MIO) algorithm. We carried out an empirical study to compare the MIO
91 algorithm with the current state-of-the-art, namely the Whole Test Suite
92 (WTS) [6] approach and the Many-Objective Sorting Algorithm (MOSA) [7].
93 We also used random search as a baseline.

94 We carried out a series of experiments on a set of artificial software with
95 different characteristics (clear gradients, plateaus, deceptive local optima and
96 infeasible targets). In most cases MIO achieves higher coverage. However,
97 artificial problems might not be fully representing the characteristics of real
98 software. Therefore, we also carried out experiments on system testing of
99 seven RESTful API web services using the open-source EVOMASTER¹ tool.

¹<http://www.evomaster.org>

Overall, MIO obtained the best results, but not on all of the seven web services. In these latter cases though, MIO still obtained the second-best results. Therefore, it is important to understand why this was the case, and what can be learned to propose novel MIO variants to improve performance even further.

To enable the replicability of this study, and to enable the use of MIO in other research and in industrial contexts, we provide the source code of all experiments and case studies as open-source on GitHub², currently the most popular open-source repository.

This article is an extension of a conference paper [8]. In particular, in this extension we analyze a further variant of how MIO can handle infeasible targets. Furthermore, the original [8] only had experiments on artificial examples and three tiny numerical functions, which are replaced here with the system testing of seven RESTful APIs.

This paper is organized as follows. Section 2 provides background information needed to understand the rest of the paper. The novel MIO algorithm is presented in Section 3. The empirical study is described in Section 4. Threats to validity are discussed in Section 5. Finally, Section 6 concludes the paper.

2. Background

2.1. Search-Based Software Testing

There has been a lot of research on how to automate the generation of high quality test cases. One of the easiest approach is to generate test cases at random [9]. Although it can be effective in some contexts, random testing is often not an effective strategy.

Among the different techniques proposed throughout the years, search-based software engineering has been particularly effective at solving many different kinds of software engineering problems [1], in particular software testing [10], with advanced tools for unit test generation like EvoSuite³ [11, 12].

Software testing can be modeled as an optimization problem, where one wants to maximize the code coverage and fault detection of the generated test suites. Then, once a fitness function is defined for a given testing problem, a search algorithm can be employed to explore the space of all possible solutions (test cases in this context).

²<https://github.com/EMResearch>

³<https://github.com/EvoSuite/evosuite>

133 There are several kinds of search algorithms, where Genetic Algorithms
 134 (GAs) are perhaps the most famous. In a GA, a population of individuals
 135 is evolved for several generations. Individuals are selected for reproduction
 136 based on their fitness value, and then go through a crossover operator (mixing
 137 the material of both parents) and mutations (small changes) when sampling
 138 new offspring. The evolution ends either when an optimal individual is
 139 evolved, or the search has run out of the allotted time.

140 2.2. Whole Test Suite (WTS)

141 The Whole Test Suite [6] approach was introduced as an algorithm to
 142 generate whole test suites. Before that, typical test case generators were
 143 targeting only single objectives, like specific lines or branches, using heuristics
 144 like the *branch distance* and the *approach level* (as for example done in [13]).

145 In the WTS approach, a GA is used, where an individual in the GA
 146 population is a set of test cases. Mutation and crossover operators can modify
 147 both the set composition (i.e., remove or add new tests) and its content (e.g.,
 148 modify the tests). As fitness function, the sum of all branch distances in the
 149 SUT is used. At the end of the search, the best solution in the population is
 150 given as output test suite. To avoid losing good tests during the search, the
 151 WTS can also be extended to use an *archive* of best tests seen so far [14].

152 2.3. Many-Objective Sorting Algorithm (MOSA)

153 The Many-Objective Sorting Algorithm (MOSA) [7] was introduced to
 154 overcome some of the limitations of WTS. In MOSA, each testing target (e.g.,
 155 lines) is an objective to optimize. MOSA is an extension of NSGA-II [15], a
 156 very popular multi-objective algorithm. In MOSA, the population is composed
 157 of tests, not test suites. When a new target is covered, the test covering it
 158 gets stored in an archive, and such target is not used any more in the fitness
 159 function. A final output test suite is composed by the best tests found during
 160 the search and that are stored in the archive.

161 In NSGA-II, selection is based on ranks (from 1 on, where 1 is the best):
 162 an individual that subsumes many other individuals gets a better rank, and so
 163 it is more likely to be selected for reproduction. One of the main differences of
 164 MOSA compared to NSGA-II is the use of the *preference sorting criterion*: to
 165 avoid losing the best individuals for a given testing target, for each uncovered
 166 testing target the best individual gets the best rank (0 in MOSA), regardless
 167 of its subsuming relations with the other tests.

168 2.4. *EvoMaster*

169 EVOMASTER⁴ is an open-source tool written in a mix of Kotlin and Java.
170 EVOMASTER aims at generating system level test cases, using search-based
171 techniques. EVOMASTER currently focuses on web services, in particular
172 RESTful APIs [16].

173 EVOMASTER is composed of two main components: a *core* process re-
174 sponsible for the main functionalities (e.g., command-line parsing, search
175 and generation of test files), and a *driver* process. This latter is responsible
176 to start/stop/reset the SUT and instrument its source code, e.g., via auto-
177 mated bytecode manipulation, in a similar way of how unit test tools like
178 EvoSuite [11] do. For example, you need to add probes in the bytecode to
179 check which statements are executed, and also to define heuristics to help
180 solving the predicates in the branch statement (e.g., the so called *branch*
181 *distance* [17]).

182 EVOMASTER generates test suites with the goal of optimising white-
183 box, code coverage metrics (e.g., statement and branch coverage) and fault
184 detection (e.g., HTTP 5xx status codes can be used in some cases as automated
185 oracles). Each test case will be composed of one or more HTTP calls. The
186 generated test files (e.g., using JUnit⁵ and RestAssured⁶ libraries) are self-
187 contained, as using the EVOMASTER driver as a library to automatically
188 start the SUT before running the tests (e.g., in JUnit this can be done in a
189 `@BeforeClass` init method).

190 Generating test cases for RESTful APIs is a complex tasks, with a very
191 large search space. Not only one needs to decide how many HTTP calls to
192 do, but for each call we also need to setup all the appropriate HTTP headers,
193 specify the HTTP verb (e.g., POST or GET), URL path parameters, URL
194 query parameters and HTTP body payloads. These latter could be arbitrarily
195 complex (many fields with string and numerical values), representing data
196 usually encoded in JSON or XML formats.

197 A REST endpoint is usually implemented with a function in a class (e.g.,
198 in Java). Server frameworks (e.g., Spring and JEE) will automatically call
199 such function when a HTTP request is received (which function is called
200 depends on the URI of the requested HTTP resource). The data in a HTTP

⁴<http://www.evomaster.org>

⁵<http://junit.org/junit4/>

⁶<https://github.com/rest-assured/rest-assured>

```

@PUT
@Timed
@Path("/{id}")
@Produces(MediaType.APPLICATION_JSON)
@UnitOfWork
public Tag update(@Auth @ApiParam(hidden = true) AuthResult authResult,
    @Context HttpServletResponse response,
    @PathParam("id") long id,
    Tag updatedTag) {
    doAuth(authResult, response, Permission.category_edit);
    Tag persisted = dao.read(id);

    persisted.setName(updatedTag.getName());
    persisted.setGroup(updatedTag.getGroup());

    dao.update(persisted);
    return persisted;
}

```

Figure 1: Example of a simple RESTful endpoint from the *scout-api* SUT.

```

@Test
public void test9() throws Exception {

    given().accept("/*/*")
        .header("Authorization", "ApiKey administrator")
        .contentType("application/json")
        .body("{\"id\":-1116301531,
            \"group\": \"MGwNpDeS8y\",
            \"media_file\": {\"mime_type\": \"xLoD2NC72Ag\",
                \"author\": \"4o9JkwgN7\"},
            \"activities_count\": -1696622830}")
        .put(baseUrlOfSut + "/api/v1/categories/-5331650344483936669")
        .then()
        .statusCode(500);
}

```

Figure 2: Example of test generated by EVOMASTER for the endpoint in Figure 1.

201 request (e.g., payload and headers) is automatically unmarshalled into the
 202 inputs of the endpoint's function (e.g., strings and data transfer objects).

203 Figure 1 shows a simple example of an endpoint from the *scout-api* SUT.
 204 In that endpoint, a PUT is handled to replace an existing resource. Such
 205 resource is first loaded from the database, and the updated values are then
 206 persisted back. Figure 2 shows an example of test generated by EVOMASTER
 207 for such endpoint, using the RestAssured library. The payload is written in
 208 JSON, which will be unmarshalled into the `Tag` Java class when the HTTP

209 request is handled by the server. Note that such test reveals a bug in the
 210 SUT, as the returned HTTP status is 500 (server error). When doing an
 211 update, the SUT does not check if the resource to update exists, leading to
 212 a null pointer exception when an invalid id is given as input. The correct
 213 behaviour in this case would had been to return a 404 status code.

214 3. The MIO Algorithm

215 3.1. Core Algorithm

216 Both WTS and MOSA have been shown to provide good results, at least
 217 for unit test generation [6, 14, 7]. However, both algorithms have intrinsic
 218 limitations, like for example:

- 219 • population-based algorithms like WTS and MOSA do put more em-
 220 phasis on the exploration of the search landscape, which is not ideal
 221 in constrained situations of limited search budgets, like for example in
 222 system-level testing where each test case execution can be computation-
 223 ally expensive. Letting the user to tune the population size parameter
 224 is not a viable option, unless it is done automatically (but even then, it
 225 has side effects, as we will see in the empirical study).
- 226 • although once a target is covered it is not used any more for the fitness
 227 function, the individuals optimised for it would still be in the population.
 228 They will die out eventually after a few generations, but, until then,
 229 their presence in the population can hamper the search if those covered
 230 targets are unrelated to the remaining non-covered targets.
- 231 • in the presence of infeasible targets, some tests can get good fitness
 232 score (e.g., a close to 0 branch distance) although they will never cover
 233 those infeasible targets. Those not useful tests might end up taking
 234 over a large part of the population.
- 235 • there can be a very large number of objectives to cover, even in the
 236 order of hundreds of thousands (e.g., in the system-level testing of
 237 industrial systems). A fixed size population would simple not work well:
 238 if too small, then there would not be enough diverse genetic material
 239 in the first generation; if too large, not only convergence would be
 240 drastically slowed down, but also the computational cost could sky-rock
 241 (e.g., NSGA-II has a quadratic complexity based on the population
 242 size).

Algorithm 1 Many Independent Objective (MIO) Algorithm

Input: Stopping condition C , Fitness function δ , Population size limit n ,
Probability of random sampling P_r , Start of focused search F

Output: Archive of optimised individuals A

```
1:  $T \leftarrow \text{SETOFEMPTYPOPULATIONS}()$ 
2:  $A \leftarrow \{\}$ 
3: while  $\neg C$  do
4:   if  $P_r > \text{rand}()$  then
5:      $p \leftarrow \text{RANDOMINDIVIDUAL}()$ 
6:   else
7:      $p \leftarrow \text{SAMPLEINDIVIDUAL}(T)$ 
8:      $p \leftarrow \text{MUTATE}(p)$ 
9:   end if
10:  for all  $k \in \text{REACHEDTARGETS}(p)$  do
11:    if  $\text{ISTARGETCOVERED}(k)$  then
12:       $\text{UPDATEARCHIVE}(A, p)$ 
13:       $T \leftarrow T \setminus \{T_k\}$ 
14:    else
15:       $T_k \leftarrow T_k \cup \{p\}$ 
16:      if  $|T_k| > n$  then
17:         $\text{REMOVETWORSTTEST}(T_k, \delta)$ 
18:      end if
19:    end if
20:  end for
21:   $\text{UPDATEPARAMETERS}(F, P_r, n)$ 
22: end while
23: return  $A$ 
```

243 To avoid these limitations, we have designed a novel evolutionary algo-
244 rithm that we call the Many Independent Objective (MIO) algorithm. In a
245 nutshell, MIO combines the simplicity and effectiveness of (1+1) EA with a
246 dynamic population, dynamic exploration/exploitation tradeoff and feedback-
247 directed target selection. Algorithm 1 provides a high level pseudo-code of
248 the MIO algorithm. The full details of the algorithm can be found (written in
249 Kotlin) in the repository of the EVOMASTER tool, in particular in the class
250 `MioAlgorithm`.

251 The MIO algorithm maintains an archive of tests. In the archive, *for*

252 *each* testing target we keep a different population of tests of size up to n (e.g.,
 253 $n = 10$). Therefore, given z objectives/targets, there can be up to $n \times z$ tests
 254 in the archive at the same time.

255 At the beginning of the search, the archive will be empty, and so a new
 256 test will be randomly generated. From the second step on, MIO will decide
 257 to either sample a new test at random (probability P_r), or will choose (details
 258 later) one existing test in the archive (probability $1 - P_r$), copy it, and mutate
 259 it. Every time a new test is sampled/mutated, its fitness is calculated, and it
 260 will be saved in the archive if needed (details later). At this point, we need
 261 to define how tests are saved in the archive, and how MIO samples from the
 262 archive.

263 When a test is evaluated, a copy of it might be saved in 0 or more of the
 264 z populations in the archive, based on its fitness value. For each target, there
 265 will be a heuristic score h in $[0,1]$, where 1 means that the target is covered,
 266 whereas 0 is the worst possible heuristic value. For example, if the heuristic
 267 is the branch distance d , this can be mapped into $[0,1]$ by using $h = 1/(1 + d)$
 268 (where $h = 0$ if a branch was never reached and so the branch distance d was
 269 not calculated).

270 For each target k , a test is saved in population T_k , with $|T_k| \leq n$, if either:

- 271 • if $h_k = 0$, the test is not added regardless of the following conditions.
- 272 • if the target is covered, i.e. $h_k = 1$, the test is added and that population
 273 is shrunk to one single individual, and it will never expand again (i.e.,
 274 it will be always $|T_k| = 1$). A new test can *replace* the one in T_k only
 275 if it is *shorter* (which will depend on the problem domain, e.g. size
 276 measured in sequence of function calls in unit testing) or, if it is of the
 277 same size, then replace the current test only if the new test has better
 278 coverage on the other targets (i.e., sum of all the heuristic values on all
 279 targets).
- 280 • if the population is not full (i.e., $|T_k| < n$), then the test is added.
 281 Otherwise, if full (i.e., $|T_k| = n$), the test might replace the worst in the
 282 population, but only if not worse than it (but not necessarily better).
 283 This means no worse heuristic value or, if the same, no larger size.

284 The idea is that, for each target, we keep a population of candidate tests
 285 for it, for which we have at least some heuristic value. But once a target k is
 286 covered, we just need to store the best test, and discard the rest. Note that,

287 if a discarded test in T_k was good for another target j , then it would be still
 288 stored in T_j anyway, so it is not lost.

289 When MIO needs to sample one test from the archive instead of generating
 290 one at random, it will do the following:

- 291 • choose one target k at random where $|T_k| > 0$ and k is not covered
 292 (i.e., no test has $h_k = 1$). If all non-empty populations are for covered
 293 targets, then just choose k randomly among them.
- 294 • choose one test randomly from T_k .

295 By using this approach, we aim at sampling tests that have non-zero
 296 heuristic (and so guidance) for targets that are not covered yet.

297 3.2. *Exploration/Exploitation Control*

298 In the MIO algorithm, the two main parameters for handling the trade-
 299 off between exploration and exploitation of the search landscape are the
 300 probability P_r of sampling at random and the population size n per target.
 301 Exploration is good at the beginning of the search, but, at the end, a more
 302 focused exploitation can bring better results. Like in Simulated Annealing,
 303 we use an approach in which we gradually reduce the amount of exploration
 304 during the search.

305 We define with F the percentage of time after which a focused search
 306 should start. This means that, for some parameters like P_r and n , we define
 307 two values: one for the start of the search (e.g., $P_r = 0.5$ and $n = 10$), and
 308 one for when the focused phase begins (i.e., $P_r = 0$ and $n = 1$). These values
 309 will linearly increase/decrease based on the passing of time. For example,
 310 if $F = 0.5$ (i.e., the focused search starts after 50% of the search budget is
 311 used), then after 30% of the search, the value P_r would decrease from 0.5 to
 312 0.2.

313 Note, when during the search decreasing n leads to some cases with
 314 $|T| > n$, then those populations are shrunk by removing the worst individuals
 315 in it. Once the focused search begins (i.e., $P_r = 0$ and $n = 1$), then MIO
 316 starts to resemble a parallel (1+1) EA.

317 When dealing with many objectives, even if there is a clear gradient to
 318 cover them in the fitness landscape, there might be simply not enough time
 319 left to cover all of them. In software testing, the final user is only interested
 320 in tests that do cover targets, and not in tests that are heuristically close to
 321 cover them (e.g., close to solve complex constraints in some branch predicates,

but not there yet). Therefore, between a test suite A that is close to but does not cover 100 targets, and another one B which does cover 1 target and is very far from covering the remaining 99, the final user would likely prefer B over A .

To take this insight into account, MIO tries to focus on just a few targets at a time, instead of spreading its resources thin among all the left uncovered targets. For example, in MIO there is an extra parameter m which controls how many mutations and fitness evaluations should be done on the same individual before sampling a new one. Like P_r and n , m varies over time, like starting from 1 and then increasing up to 10 when the focused search begins.

3.3. Feedback-Directed Sampling

When dealing with many objectives and limited resources, it might not be possible to cover all of them. As discussed in Section 3.2, the final user is only interested in the actually covered targets, and not on how close we are to cover them. Therefore, it makes sense to try to focus on targets that we have higher chances to cover. This is helpful also when dealing with infeasible targets for which any heuristic will just plateau at a certain point.

To handle these cases, we use a simple but yet effective technique that we call Feedback-Directed Sampling (FDS). The sampling algorithm from the archive discussed in Section 3.1 is modified as follow. Instead of choosing the target k randomly among the non-covered/non-empty ones, each of these targets will have a counter c_k . Every time a test is sampled from a population T_k , then c_k is increased by 1. Every time a new *better* individual is added to T_k (or replace one of its existing tests), then the counter c_k is reset to 0. When we sample from k from non-covered/non-empty ones, then, instead of choosing k at random, we choose the k with the lowest c_k .

As long as we get improvements for a target k , the higher chances will be that we sample from T_k , as c_k gets reset more often. On the other hand, for infeasible targets, their c will never be reset once they reach their plateau, and so they will be sampled less often. Similarly, more complex targets will be sampled less often, and so the search concentrates on the easier targets that are not covered yet. However, this is not an issue because, once an easy target k is covered, we do not sample from T_k any more (recall Section 3.1), unless also *all* the other targets are either covered or with empty T .

However, there is a potential issue with this kind of FDS based on *last* (most recent) improvement. For example, consider the case of d difficult, albeit feasible targets, for which it might take many fitness evaluations

before getting an improvement. Their c counters will have same/similar values. However, instead of focusing in solving at least one of these feasible d targets within the given budget constraints, the FDS would spread resources equally among them. In other words, as soon as a difficult target j gets a c_j greater than the others, it will take at least $d - 1$ fitness evaluations before sampling again from the population T_j . In those cases, it would make sense to concentrate resources to cover at least one of these d targets, instead of spreading such resources equally among these d targets and likely cover none of them.

A possible approach is to use a more *focused* FDS. For each target, we also keep track of how many steps s it took to get the last improvement in fitness for that target. In other words, when a target j gets a fitness improvement, we save $s_j = c_j$ before resetting $c_j = 0$. When choosing a T to sample from based on c counters, instead of choosing the target j based on the lowest c , we keep sampling from the previously sampled j , but only as long as $c \leq 2s$. In other words, once we get an improvement in fitness within s steps for a target j , we keep sampling from T_j for a number of times that is no more than twice it took to get to the previous improvement. Once the condition $c_j \leq 2s_j$ does not hold any longer for the current target j , we choose a different target to sample from with the *last* FDS approach.

Whether FDS improves or not performance is a matter of empirical investigation, as it depends on the proportion and type of infeasible targets in the SUTs. As such, FDS is not a core component of MIO, and rather a feature that can be (de)activated on demand.

4. Empirical Study

To evaluate the performance of the MIO algorithm, we compared it with random search, MOSA and WTS. We used two different case studies: (1) a set of artificial problems with varying, specific characteristics; (2) seven RESTful API web services.

In this paper, we aim at answering the following research questions:

RQ1: On which kinds of problem does MIO perform better than Random, MOSA and WTS?

RQ2: What is the impact of tuning parameters for exploration vs. exploitation of the search landscape in MIO and MOSA?

RQ3: How do the analysed algorithms fare on actual software?

394 4.1. Artificial Software

395 In this paper, we designed four different kinds of artificial problems. In
 396 all of them, there are z targets, and the search algorithm can be run for
 397 up to b fitness evaluations. A test is defined by two components: an *id*
 398 (e.g., think about it like the name of a method to call in unit testing) and
 399 a numeric integer value $x \in [0, r]$ (e.g., think about it like the input to a
 400 method call). Each target k is independent, and can be covered only by a
 401 test with $id = k$. The artificial problems will differ based on their fitness
 402 landscape. Given $g \in [0, r]$ the single global optimum chosen at random, and
 403 given the normalising function $\rho(d) = 1/(1 + d)$ for distances, then we have
 404 four different cases for each target:

405 **Gradient:** $h_k = \rho(|x - g|)$. This represents the simplest case where the search
 406 algorithm has a direct gradient from x toward the global optimum g .

407 **Plateau:** $h_k = \rho(g - x)$ if $g \geq x$, else $h_k = \rho(0.1 \times r)$. In this case, we have
 408 one side of the search landscape (before the value of the global optimum
 409 g) with a clear gradient. However, the other side is a plateau with a
 410 relatively good fitness value (note that $0 \leq |g - x| \leq r$).

411 **Deceptive:** $h_k = \rho(g - x)$ if $g \geq x$, else $h_k = \rho(1 + r - x)$. This is similar
 412 to the *Plateau* case, where one side of the search landscape has a clear
 413 gradient toward g . However, the other side has a deceptive gradient
 414 toward leaving g and reach the maximum value r .

415 **Infeasible:** like *Gradient*, but with a certain number of the z targets having
 416 a constant $h_k = \rho(1)$ and no global optimum.

417 We implemented the four search algorithms in which, when a test is
 418 sampled, its *id* and x values are chosen at random within the given valid
 419 ranges. Mutations on x is done by adding/subtracting 2^i , where i is chosen
 420 randomly in $[0, 10]$. We consider mutating *id* as a *disruptive* operation, and,
 421 as such, we only mutate it with low probability 0.01. Mutating *id* means
 422 changing both *id* and x at random (think about mutating a function call
 423 with string inputs into another one that requires integers, where the strings
 424 x would have no meaning as integers). All the analysed search algorithms
 425 use the same random sampling, mutation operation and archive to store the
 426 best tests found so far.

427 For the MIO algorithm, we used $F = 0.5$, $P_r = 0.5$, $n = 10$ and max
 428 mutations 10. For MOSA, we used the same settings as in [7], i.e. population




Figure 3: Coverage results on the *Gradient* problem type, with varying number of targets z .

size 50 and tournament selection size 10. WTS uses the same population size as MOSA, with up to 50 test cases in the same test suite (i.e., one individual). A randomly sampled test suite in WTS will have size randomly chosen between 1 and 50. WTS also has mutation operators to add a new test (probability 1/3) in a test suite, remove one test at random (probability 1/3), or modify one (probability 1/3) like in MIO and MOSA. WTS also uses a crossover operator with probability 70% to combine test suites.

For each problem type but *Infeasible*, we created problems having a variable number of z targets, in particular $z \in \{1, 2, 3, 4, 5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100\}$, i.e., 15 different values in total, ranging from 1 to 100. We used $r = 1000$. We ran each of the four search algorithms 100 times with budget $b = 1000$. As the optima g are randomised, we make sure that the search algorithms run on the same problem instances. In the case of the *Infeasible* type, we used 10 *Gradient* targets, on which we added a different number of infeasible targets in $\{0, 1, 2, 3, 4, 5, 10, 20, 30, 40, 50, 60, 70, 80, 90, 100\}$, i.e., 16 values in total, with z ranging from $(10 + 0) = 10$ to $(10 + 100) = 110$. Figure 3 shows the results for the *Gradient* type, Figure 4 for *Plateau*, Figure 5 for *Deceptive*, and Figure 6 for *Infeasible*.

The *Gradient* case (Figure 3) is the simplest, where the four search




Figure 4: Coverage results on the *Plateau* problem type, with varying number of targets z .

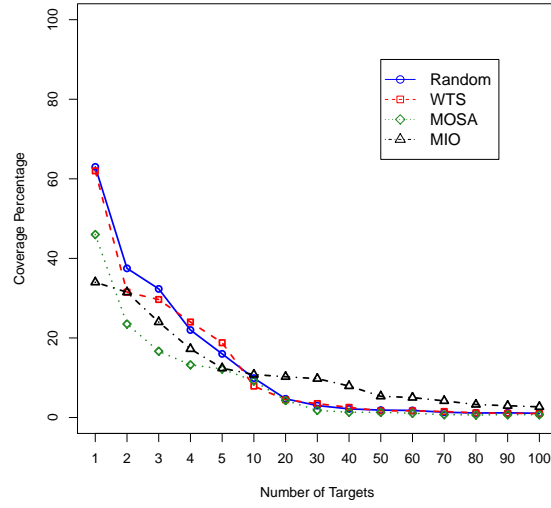


Figure 5: Coverage results on the *Deceptive* problem type, with varying number of targets z .




Figure 6: Coverage results on the *Infeasible* problem type, with varying number of infeasible targets on top of 10 *Gradient* ones.

448 algorithms obtain their highest coverage. MIO and MOSA have very similar
 449 performance, which is higher than the one of Random and WTS. However,
 450 on the more difficult case of *Plateau* (Figure 4), MIO starts to have a clear
 451 advantage over MOSA. For example, from $z = 30$ on, MOSA becomes
 452 equivalent to Random and WTS, covering nearly no target. However, in that
 453 particular case, MIO can still achieve around 20% coverage (i.e., 6 targets).
 454 Even for large numbers of targets (i.e., 100 when taking into account that
 455 the search budget b is only 1000), still MIO can cover some targets, whereas
 456 the other algorithms do not.

457 The *Deceptive* case (Figure 5) is of particular interest: for low numbers of
 458 z targets (i.e., up to 10), both MIO and MOSA perform worse than Random.
 459 From 10 targets on, MOSA is equivalent to Random and WTS, whereas
 460 MIO has better results. This can be explained by taking into account two
 461 contrasting factors: (1) the more emphasis of MIO and MOSA on exploitation
 462 compared to the exploration of the search landscape is not beneficial in
 463 deceptive landscape areas, whereas a random search would not be affected by
 464 it; (2) MIO does better handle large numbers of targets (Figure 3), even when
 465 there is no gradient (Figure 4). The value $z = 10$ seems to be the turning
 466 point where (2) starts to have more weight than (1).

467 The *Infeasible* case (Figure 6) is where MIO obtains the best results
 468 compared to the other algorithms. For this case, we also ran a further version
 469 of MIO in which we deactivated FDS (recall Section 3.3), as we wanted to
 470 study its impact in the presence of infeasible targets. From 20 infeasible
 471 targets on, MOSA, Random and WTS become equivalent, covering nearly no
 472 target. However, MIO can still cover nearly 80% of the 10 feasible testing
 473 targets. For very large numbers of infeasible targets like 100, still MIO
 474 can cover nearly 40% of the feasible ones. This much better performance is
 475 mainly due to the use of FDS (see the gap in Figure 6 between MIO and
 476 MIO-noFDS). However, even without FDS, MIO still does achieve better
 477 results compared to the other algorithms.



478 **RQ1:** *on all the considered problems, MIO is the algorithm that scaled
 best. Coverage improvements were even up to 80% in some cases.*

479 When using a search algorithm, some parameters need to be set, like the
 480 population size or crossover probability in a GA. Usually, common settings
 481 in the literature can already achieve good results on average [18]. Finding
 482 tuned settings that work better on average on a large number of different
 483 artefacts is not trivial. Ideally, a user should just choose for how long a
 484 search algorithm should run, and not do long tuning phases by himself on
 485 his problem artefacts. Parameter tuning can also play a role in algorithm
 486 comparisons: what if a compared algorithm performed worse just because
 487 one of its chosen settings was sub-optimal?



488 Arguably, among the most important parameters for a search algorithm
 489 are the ones that most impact the tradeoff between the exploration and
 490 the exploitation of the search landscape. In the case of MIO, this is clearly
 491 controlled by the F parameter (low values put more emphasis on exploitation,
 492 whereas for high values a large number of tests are simply sampled at ran-
 493 dom). In the case of population-based algorithms, the population size can be
 494 considered as a parameter to control such tradeoff. Small populations would
 495 reward exploitation, whereas large populations would reward exploration.

496 To study these effects, we carried out a further series of experiments on
 497 the *Gradient*, *Plateau* and *Deceptive* problem types. For MIO, we studied
 498 six different values for F , in particular $\{0, 0.2, 0.4, 0.6, 0.8, 1\}$. For MOSA, we
 499 studied six different values for the population size, i.e. $\{4, 8, 16, 32, 64, 128\}$.
 500 Each experiment was repeated 100 times. Figure 7 shows the results of these
 501 experiments.



502 For MIO, the results in Figure 7 do match expectation: for problems



(a) *Gradient* problem type.



(b) *Plateau* problem type.



(c) *Deceptive* problem type.

Figure 7: Tuning of F for MIO (left side) and population size for MOSA (right side).

503 with clear gradient or with just some plateaus, a more focused search that
 504 rewards exploitation is better. The best setting is a low $F = 0.2$, although the
 505 lowest $F = 0$ is not particularly good. You still need some genetic diversity
 506 at the beginning of the search, and not rely on just one single individual. For
 507 deceptive landscapes, exploration can be better, especially for a low number of
 508 targets. For example, with $z = 1$ then $F = 1$ provides the best performance.
 509 However, for larger number of targets, too much exploration would not be so
 510 beneficial, as it would not have enough time to converge to cover the targets.

511 In the case of MOSA, Figure 7 provides some interesting insight. For
 512 simple problems with clear gradient, one would expect that a focused search
 513 should provide better results. However, the small population size of 4 is
 514 actually the configuration that gave the worst results. The reason is that
 515 there is only little genetic material at the beginning of the search, and new
 516 one is only generated with the mutation operator. However, a too large
 517 population size would still be detrimental, as not focused enough. In that
 518 particular problem type, the best population size seems to be ranging from
 519 16 to 32, i.e., not too large, but not too small either. In case of plateaus,
 520 still a too small population size (e.g., 4) gives the worst result. However,
 521 in case of plateaus, there is a need to have some more exploration in the
 522 search landscape, and this is confirmed by the fact that the best results are
 523 obtained with large population sizes (e.g., 64 and 128). This effect is much
 524 more marked in the case of deceptive landscapes, where large population sizes
 525 lead to much better results.

526 The experiments reported in Figure 7 clearly points out to a challenge in
 527 population-based algorithms when dealing with many-objective problems. A
 528 too small population size would reduce diversity in the initial genetic material.
 529 But a too large population size would hamper convergence speed. Finding a
 530 fixed, right population size that works on most problem sizes (e.g., $z = 10$
 531 vs $z = 1,000,000$) might not be feasible. To overcome this issue, MIO uses
 532 a dynamically sized population, whereas the tradeoff between exploration
 533 and exploitation is controlled by a dynamically decreasing probability P_r of
 534 creating new tests at random (instead of mutating the current ones stored in
 535 the archive).

536 **RQ2:** *On the analysed problems, the population size and the F parameter have clear effects on performance, which strongly depend on whether on the given problem one needs more or less exploitation/exploration.*

Table 1: Information about the seven RESTful web services used in the empirical study. We report their number of Java/Kotlin classes and lines of code. We also specify the number of endpoints, i.e., the number of exposed resources and HTTP methods applicable on them, and also whether these SUTs interact with a database.

Name	# Classes	LOCs	Endpoints	Database
<i>catwatch</i>	69	5442	23	yes
<i>features-service</i>	23	1247	18	yes
<i>proxyprint</i>	68	7534	115	yes
<i>rest-ncs</i>	9	602	6	no
<i>rest-news</i>	10	715	14	yes
<i>rest-scs</i>	13	859	11	no
<i>scout-api</i>	75	7479	49	yes
Total	267	23878	236	5/7 yes

537 4.2. Real Software

538 When designing algorithms to work on a large class of problems, it is
539 common to evaluate them on artificial problems to try to abstract away and
540 analyse in details the characteristics for which such algorithms perform best.
541 For example, the very popular NSGA-II algorithm (on which MOSA is based
542 on) was originally evaluated only on nine numerical functions [15]. However,
543 using only artificial problems is risky, as those might abstract away some very
544 important factors. A good example of this issue is Adaptive Random Testing,
545 where artificial problems with artificially high fault rates were masking away
546 its prohibitive computational cost [19].

547 To mitigate this issue, we also carried out experiments on actual software,
548 using the EVOMASTER tool, where we implemented the MIO algorithm
549 besides WTS and MOSA. EVOMASTER generates system level test cases for
550 RESTful API web services. Seven APIs were chosen for the experiments, and
551 they are described in Table 1. To be able to replicate the experiments in this
552 paper, we provide this case study online on GitHub⁷.

553 The APIs *features-service*, *proxyprint*, *scout-api* and *catwatch* are actual
554 open-source Java projects selected from GitHub, using different technologies
555 like Spring and Dropwizard. These four APIs use a SQL database. On

⁷<https://github.com/EMResearch/EMB>

the other hand, *rest-news* is an artificial SpringBoot application written in Kotlin, which has been used in didactic settings. Similarly, *rest-ncs* (REST Numerical Case Study) and *rest-scs* (REST String Case Study) are artificial APIs using SpringBoot and that contain code previously used for experiments in unit testing of classes heavily dependent on numerical [19] and string [20] computations.

We ran MIO, MOSA, WTS and random search on each of the seven web services. For MIO, we considered three variants: without FDS, *last* FDS and *focused* FDS. Each algorithm was run with the same search budget of 100 thousand HTTP calls. Each experiment was repeated 50 times with different random seeds. Each experiments required at least three CPUs, as three processes needed to be run in parallel and the SUTs are multi-threaded. Therefore, a large cluster of computers was required to run all these experiments, which needed a total of 632 days of CPU time.

When comparing algorithms, it is also important to study their computational cost. An algorithm could be more efficient, and run more fitness evaluations given the same amount of time. But the computational cost is strongly affected by the low level details of the code implementation, and so it poses further threats to validity for fair comparisons. We did not consider the computational cost in this paper because, for system testing, the cost of a fitness evaluation (marshalling of JSON data, HTTP over TCP and access to SQL databases) is much higher than any overhead of the used algorithm. This, however, might not be the case in the context on unit testing, where using time as stopping criterion would be more appropriate (as such algorithm overheads might not be negligible compared to the cost of the fitness function).

Average values of these experiments are reported in Table 2, where we also report the Vargha-Delaney effect sizes \hat{A}_{12} and the results of Mann-Whitney-Wilcoxon U-tests at $\alpha = 0.05$ level [21]. Table 3 reports the relative rank in performance among the different algorithms, which is statistically significant according to the Friedman test. Note that, for MIO, in these tables we only report the *focused* FDS variant.

From what can be seen in Table 2 and Table 3, the MIO algorithm has overall the best results. However, it is not always the best, i.e., only on three out of seven SUTs. On the other hand, on these other SUTs MIO is the second best algorithm. Random search is generally the worst algorithm. MOSA and WTS have inconsistent results. For example, although MOSA is the best on *rest-ncs* and *rest-scs*, it is also the worst on *proxyprint*, even

Table 2: Comparisons of algorithms on the seven web services. Coverage is not a percentage, but rather the average raw sum of covered targets used in the fitness function of EVOMASTER. For each algorithm, we also specify if better than any of the others, i.e. $\hat{A}_{12} > 0.5$ (in parenthesis) and p-value less than 0.05.

SUT	Algorithm	Tests	Coverage	Better than
<i>catwatch</i>	MIO	51.2	1000.0	MOSA(0.64) RAND(0.81) WTS(0.75)
	MOSA	52.3	991.3	RAND(0.72) WTS(0.63)
	RAND	52.5	985.6	
	WTS	49.9	990.6	RAND(0.64)
<i>features-service</i>	MIO	51.6	558.8	RAND(0.78)
	MOSA	52.2	551.4	RAND(0.76)
	RAND	38.1	479.1	
	WTS	63.0	592.4	MIO(0.66) MOSA(0.69) RAND(0.87)
<i>proxyprint</i>	MIO	258.0	1559.8	MOSA(0.95) RAND(0.85) WTS(0.84)
	MOSA	298.3	1526.5	
	RAND	300.4	1534.1	MOSA(0.63)
	WTS	294.5	1538.9	MOSA(0.74)
<i>rest-ncs</i>	MIO	55.2	525.6	RAND(1.00) WTS(0.86)
	MOSA	51.9	526.5	RAND(1.00) WTS(0.90)
	RAND	41.9	436.4	
	WTS	48.4	515.9	RAND(1.00)
<i>rest-news</i>	MIO	30.3	265.0	
	MOSA	31.1	265.3	RAND(0.63)
	RAND	31.1	262.7	
	WTS	31.8	265.3	
<i>rest-scs</i>	MIO	73.5	594.1	RAND(1.00) WTS(0.98)
	MOSA	109.6	722.0	MIO(1.00) RAND(1.00) WTS(1.00)
	RAND	34.3	510.2	
	WTS	53.1	534.4	RAND(0.99)
<i>scout-api</i>	MIO	173.2	1829.5	RAND(1.00)
	MOSA	188.4	1828.0	RAND(1.00)
	RAND	177.1	1427.8	
	WTS	200.9	1944.0	MIO(0.99) MOSA(0.99) RAND(1.00)

594 worse than random search. Similarly, WTS is the best on *features-service*
595 and *scout-api*, but then it is the second-worst on four of the remaining SUTs.

596 To shed more light on these different performance behaviours, Figure 8
597 plots the average number of covered targets through time. The behaviour




Figure 8: Performance over time of the different search algorithms.

Table 3: Algorithm ranks based on their average performance, where Rank 1 represents the highest achieved coverage. We also report the χ^2 and p -value of the Friedman test.

SUT	MIO	MOSA	WTS	RAND
<i>catwatch</i>	1	2	3	4
<i>features-service</i>	2	3	1	4
<i>proxyprint</i>	1	4	2	3
<i>rest-ncs</i>	2	1	3	4
<i>rest-news</i>	1	2	3	4
<i>rest-scs</i>	2	1	3	4
<i>scout-api</i>	2	3	1	4
Average	1.6	2.3	2.3	3.9
$\chi^2=11.74286$, p -value=0.008317995				

of MOSA is of particular interest. During its initial phase, it is worse than random search on *catwatch*, *rest-news* and *scout-api*. It takes it 20%-25% of the search budget before becoming better than random search. However, on *rest-ncs* and *rest-scs* it has a clear better margin already from the very beginning, i.e. in the first 5% of the budget. On *rest-scs* in particular, MOSA has a very large margin over all the other algorithms.

At the current moment, we are unable to explain such peculiar behaviour. But, to design better algorithms, it will be important to study in detail why this kind of performance improvements are obtained. However, to enable researchers to get such kind of insight, it will be necessary to develop tools to visualize the evolution of these algorithms through time. For example, if such kind of behaviour is linked to some specific properties of the SUTs, and if such properties can be detected before starting a search, a hybrid approach could be designed to choose the best algorithm (e.g., MIO or MOSA) to employ based on those detected characteristics.

To address the potential issue of infeasible targets, the MIO algorithm employs the FDS technique, with two different variants: *last* and *focused*. Comparisons of these variants of FDS are shown in Table 4. On artificial software, FDS gives very large benefits (e.g., recall Figure 6). However, on actual software, improvements are relatively modest, if any at all. In two cases (*proxyprint* and *rest-news*), *focused* FDS gives statistically better results. However, on two other cases (*rest-scs* and *scout-api*), FDS actually

Table 4: Comparisons of MIO FDS variants on the seven web services. Coverage is not a percentage, but rather the average raw sum of covered targets used in the fitness function of EVOMASTER. For each variant, we also specify if better than any of the others, i.e. $\hat{A}_{12} > 0.5$ (in parenthesis) and p-value less than 0.05.

SUT	FDS	Tests	Coverage	Better than
<i>catwatch</i>	NONE	51.4	1003.7	
	LAST	51.2	1000.0	
	FOCUSED	51.3	998.8	
<i>features-service</i>	NONE	51.5	552.4	
	LAST	51.6	558.8	
	FOCUSED	50.8	555.9	
<i>proxyprint</i>	NONE	268.5	1546.6	
	LAST	258.0	1559.8	NONE(0.70)
	FOCUSED	261.2	1566.3	NONE(0.73)
<i>rest-ncs</i>	NONE	55.3	526.5	
	LAST	55.2	525.6	
	FOCUSED	54.8	525.1	
<i>rest-news</i>	NONE	30.0	267.2	
	LAST	30.3	265.0	
	FOCUSED	30.5	270.0	LAST(0.62)
<i>rest-scs</i>	NONE	79.6	609.7	LAST(0.62)
	LAST	73.5	594.1	
	FOCUSED	73.3	597.2	
<i>scout-api</i>	NONE	177.7	1883.0	LAST(0.71) FOCUSED(0.69)
	LAST	173.2	1829.5	
	FOCUSED	174.7	1844.3	

620 gives statistically *worse* results.

621 A possible explanation is that, in the artificial examples, all targets
622 were completely independent. However, on actual software, there can be
623 dependencies. For example, a test case whose execution does reach an
624 infeasible target can also reach other targets. Consider this simple example:

```
625 void foo(int x, int y){
626     if(complexPredicate(x)){
```




Figure 9: Coverage results in the example of high reachability among targets, with increasing number of infeasible targets out of 100.

```

627
628     if(infeasiblePredicate(y)){
629         // target i
630     }
631     if(feasiblePredicate(y)){
632         // target j
633     }
634 }
635 }

```

636 The target i is infeasible. However, any test that does reach the predicate
637 for i does also necessarily reach the predicate for target j . Therefore, every
638 time we sample from the population for the infeasible i we still have tests that
639 could be good for other targets as well. In this particular example, sampling
640 from i would still give you the right input x value to reach target j . However,
641 this does not fully explain why in some cases the use of FDS does reduce
642 performance.

643 To provide more insight, we carried out an experiment to study how FDS
644 would behave in such type of context. We created an artificial example to
645 represent maximum dependency among targets reachability (each target is
646 always reached, albeit not necessarily covered). In this case, optimised tests
647 for infeasible targets are still useful, because they always reach as well all the

648 other feasible targets that are not covered yet.

649 We considered a function with two inputs, x and y . There are 100
650 testing targets, 50 in the form `if(x == ?)`, and 50 in the form `if(y == ?)`.
651 These branches are all at the same level (no nesting), and contain no return
652 statements. Therefore, when an input pair x, y is given as input, each single
653 target is reached, and its predicate is evaluated.

654 The constants in the `if` statements are sampled between 0 and 10,000.
655 The algorithms do not generate values out of this range. We considered cases
656 in which we inject a certain number of infeasible targets for the x branches, in
657 which the constant is -1 (the target is so not reachable, and the algorithms
658 have gradient to decrease x down till 0). In particular, we use the following
659 number of infeasible targets: $\{1, 2, 3, 4, 5, 10, 20, 30, 40, 50\}$. We ran MIO in three
660 modes: without FDS, and with *last* and *focused* FDS modes. Search budget
661 was 10,000 fitness evaluations. Experiments were repeated 100 times.

662 Figure 9 shows the results of this experiment. On one hand, for low
663 numbers of infeasible targets, FDS gives slightly worse results. On the other
664 hand, for high numbers of infeasible targets, FDS gives better results.

665 This example is deliberately constructed to represent a case in which
666 FDS would not be expected to be particularly useful, even in the presence
667 of a large proportion of infeasible targets. So, it is not unexpected that,
668 even with 50% of infeasible targets, FDS only provides minor improvement.
669 However, for small numbers of infeasible targets (up to 10%), FDS actually
670 gives slightly worse results, similarly to some of the experiments on actual
671 software. It seems like that the use of FDS does have impact on the tradeoff
672 between exploration and exploitation of the search landscape. Depending on
673 the search landscape, this can improve or worsen the performance. For small
674 number of infeasible targets, the impact of such difference seems larger than
675 the benefits of handling the infeasible targets. However, detailed control flow
676 analyses of these cases and landscape analyses will be required to shed more
677 light on this kind of dynamics.

678 **RQ3:** *the experiments on actual software are consistent with the ones on
artificial problems: the MIO algorithm still achieves the best results, but not
on all problems.*

679 5. Threats to Validity

680 Threats to *internal* validity come from the fact that our study is based
681 on the EVOMASTER tool. We cannot guarantee that it is bug free. However,

682 to reduce such threat, not only EVOMASTER has been carefully tested, but
683 also its source code is available for review as an open-source project hosted
684 on GitHub.

685 It is possible that, when comparing a new algorithm with the existing state-
686 of-the-art, better performance might be just due to wrong implementation of
687 such existing algorithms due to misunderstanding of their details. In the case
688 of WTS, we were among its original co-authors [6] . In the case of MOSA,
689 our implementation was reviewed and fixed by one of its co-authors⁸.

690 Randomized algorithms are affected by chance. To keep this factor under
691 control, each experiment was repeated either 50 or 100 times. The appropriate
692 statistical tests and effect sizes were then employed to analyze the results.

693 Threats to *external* validity come from the fact that only seven web services
694 for a total of 23 thousand lines of code were used for the empirical evaluation
695 on real software. This was due to the large cost of experiments on system
696 level test generation, which required more than 600 days of computational
697 effort. Furthermore, albeit very common in industry, RESTful APIs are not
698 so common among open-source projects. The selected sample is biased in
699 regard of what could be found and compiled/run without problems.

700 6. Conclusion

701 In this paper, we have presented a novel search algorithm that is tailored
702 for the problem of generating test suites, in particular for system testing. We
703 call it the Many Independent Objective (MIO) algorithm. We have carried
704 out an empirical study to compare MIO with the other main algorithms
705 for test suite generation: the Whole Test Suite (WTS) approach and the
706 Many-Objective Sorting Algorithm (MOSA). We also used random search as
707 a baseline.

708 On artificial problems with increasing complexity, MIO achieved better
709 results than the other algorithms. Such improvements were also confirmed
710 when using MIO for test suite generation for system testing of seven RESTful
711 API web services. However, MIO did not give the best result on all of these
712 web services.

713 Future work will focus on analyzing and understanding the reasons why,
714 in some cases, MOSA and WTS gave better results than MIO. This will be

⁸<https://github.com/EMResearch/EvoMaster/pull/1>

715 essential to design novel MIO variants to improve performance even further.
716 However, system level test suite generation is a very complex task to analyze
717 and understand. Therefore, one essential task will be to implement tools to
718 visualize the evolution of test suites throughout the search, with the aim of
719 helping researchers to analyze those complex dynamics.

720 To help researchers integrate MIO in their frameworks, all the code used
721 for the experiments in this paper is available online on a public repository, as
722 part of the EVOMASTER tool at www.evomaster.org.

723 **Acknowledgments.** We would like to thank Annibale Panichella for
724 his help with the implementation of the MOSA algorithm. This work is
725 supported by the Research Council of Norway (project on Evolutionary
726 Enterprise Testing, 274385), and by the National Research Fund, Luxembourg
727 (FNR/P10/03).

- 728 [1] M. Harman, S. A. Mansouri, Y. Zhang, Search-based software engineering:
729 Trends, techniques and applications, ACM Computing Surveys (CSUR)
730 45 (1) (2012) 11.
- 731 [2] D. H. Wolpert, W. G. Macready, No free lunch theorems for optimization,
732 IEEE Transactions on Evolutionary Computation 1 (1) (1997) 67--82.
- 733 [3] S. Ali, M. Z. Iqbal, A. Arcuri, L. C. Briand, Generating test data from
734 OCL constraints with search techniques, IEEE Transactions on Software
735 Engineering (TSE) 39 (10) (2013) 1376--1402.
- 736 [4] G. Fraser, A. Arcuri, Achieving scalable mutation-based generation of
737 whole test suites, Empirical Software Engineering (EMSE) 20 (3) (2015)
738 783--812.
- 739 [5] B. Li, J. Li, K. Tang, X. Yao, Many-objective evolutionary algorithms:
740 A survey, ACM Computing Surveys (CSUR) 48 (1) (2015) 13.
- 741 [6] G. Fraser, A. Arcuri, Whole test suite generation, IEEE Transactions
742 on Software Engineering 39 (2) (2013) 276--291.
- 743 [7] A. Panichella, F. Kifetew, P. Tonella, Automated test case generation
744 as a many-objective optimisation problem with dynamic selection of the
745 targets, IEEE Transactions on Software Engineering (TSE) 44 (2) (2018)
746 122--158.

- 747 [8] A. Arcuri, Many Independent Objective (MIO) Algorithm for Test Suite
748 Generation, in: International Symposium on Search Based Software
749 Engineering (SSBSE), 2017, pp. 3--17.
- 750 [9] A. Arcuri, M. Z. Iqbal, L. Briand, Random testing: Theoretical results
751 and practical implications, IEEE Transactions on Software Engineering
752 (TSE) 38 (2) (2012) 258--277.
- 753 [10] S. Ali, L. Briand, H. Hemmati, R. Panesar-Walawege, A systematic
754 review of the application and empirical investigation of search-based
755 test-case generation, IEEE Transactions on Software Engineering (TSE)
756 36 (6) (2010) 742--762.
- 757 [11] G. Fraser, A. Arcuri, EvoSuite: automatic test suite generation for
758 object-oriented software, in: ACM Symposium on the Foundations of
759 Software Engineering (FSE), 2011, pp. 416--419.
- 760 [12] G. Fraser, A. Arcuri, A large-scale evaluation of automated unit test
761 generation using EvoSuite, ACM Transactions on Software Engineering
762 and Methodology (TOSEM) 24 (2) (2014) 8.
- 763 [13] M. Harman, P. McMinn., A theoretical and empirical study of search
764 based testing: Local, global and hybrid search., IEEE Transactions on
765 Software Engineering (TSE) 36 (2) (2010) 226--247.
- 766 [14] J. M. Rojas, M. Vivanti, A. Arcuri, G. Fraser, A detailed investigation
767 of the effectiveness of whole test suite generation, Empirical Software
768 Engineering (EMSE) (2016) 1--42.
- 769 [15] K. Deb, A. Pratap, S. Agarwal, T. Meyarivan, A fast and elitist multiob-
770 jective genetic algorithm: NSGA-II, IEEE Transactions on Evolutionary
771 Computation (TEVC) 6 (2) (2002) 182--197.
- 772 [16] A. Arcuri, RESTful API Automated Test Case Generation, in: IEEE
773 International Conference on Software Quality, Reliability and Security
774 (QRS), IEEE, 2017, pp. 9--20.
- 775 [17] P. McMinn, Search-based software test data generation: A survey,
776 Software Testing, Verification and Reliability 14 (2) (2004) 105--156.

- 777 [18] A. Arcuri, G. Fraser, Parameter tuning or default values? an empirical
778 investigation in search-based software engineering, *Empirical Software*
779 *Engineering (EMSE)* 18 (3) (2013) 594--623.
- 780 [19] A. Arcuri, L. Briand, Adaptive random testing: An illusion of effec-
781 tiveness?, in: *ACM Int. Symposium on Software Testing and Analysis*
782 *(ISSTA)*, 2011, pp. 265--275.
- 783 [20] M. Alshraideh, L. Bottaci, Search-based software test data generation
784 for string data using program-specific search operators, *Software Testing,*
785 *Verification, and Reliability* 16 (3) (2006) 175--203. doi:[http://dx.](http://dx.doi.org/10.1002/stvr.v16:3)
786 [doi.org/10.1002/stvr.v16:3](http://dx.doi.org/10.1002/stvr.v16:3).
- 787 [21] A. Arcuri, L. Briand, A hitchhiker's guide to statistical tests for assess-
788 ing randomized algorithms in software engineering, *Software Testing,*
789 *Verification and Reliability (STVR)* 24 (3) (2014) 219--250.